Ballarch: Boundaries in Time and Honest Outcomes, from Register-Transfer Logic to Services

Vladislav Tomilov
Independent researcher

Ballarch repository: `https://github.com/4wl2d/Pokeball`

October 2026
Specification version 2.0.0-rc.1

Abstract

Software architecture styles draw boundaries in space: they constrain which code may depend on which. They say little about *when* an effect may leave a component, or what it may claim after a timeout, when retrying can charge a customer twice and reporting failure can hide a charge that happened.

We present Ballarch, which adds two *boundaries in time*: the *commit horizon*, the atomic commit of the decision that requests an effect, and the *send horizon*, the first moment the request may have left the runtime. Transactions, transports and durable execution use these instants locally; Ballarch makes them boundaries of every interaction. Each piece of state has one owner that decides in pure functions. A runtime commits state and requested effects atomically, and ends every outbound request in exactly one outcome, Done, NotDone or Unknown, chosen from its phase and its port's effect class. For operations made of several requests, a knowledge algebra over Kleene's strong three-valued logic bounds what a component may claim. It is sound, and exact for independent calls; a composite that replies with settled knowledge is an honest component one level up.

The rules are stated once. Level-specific mechanisms check the call-protocol rules in a Kotlin runtime (trace oracles, an Alloy model), C firmware (bounded power-loss exploration) and a SystemVerilog block (netlist checks and a fault-injecting bench). In a benchmark of one checkout feature written in six styles and run 2,000 times per fault setting, idiomatic layered, actor, MVU and outbox-plus-saga versions violated a correctness requirement in 12–73% of runs. Ballarch and a layered version hardened by hand violated none. Against the hardened version, Ballarch used about 29% more code and many more framework types, its classes scored lower on cohesion, and its code had higher cyclomatic complexity and a larger distance from the main sequence. It had the lowest propagation cost of the six designs and placed 88% of its code in pure classes.

Keywords: software architecture; fault tolerance; idempotency; three-valued logic; output commit; hardware/software co-design.

1 Introduction

A payment provider that does not deduplicate requests receives a charge request and charges the card; the answer is then lost. The client times out. What should it record?

- If it retries, the customer may be charged twice.
- If it reports a failure, a charged customer is told that nothing happened.

The only honest answer is a third one: the charge *may* have happened, and someone has to find out. The "third outcome" of remote procedure calls is old knowledge [8]. Distributed-systems testing records such operations as indeterminate [37, 38], and transaction standards report heuristic and hazard outcomes [77]. HTTP advises a client not to retry a non-idempotent request automatically unless it knows that the request is in fact idempotent or that it was never applied; HTTP/2 lets a server supply evidence of the latter, and gRPC retries transparently on such evidence [20, 23, 73]. Durable-execution engines record when each activity is scheduled and when it is dispatched [70]. Hardware solved a related problem long ago: processors retire instructions in order so that exceptions are precise [66], and memory types that model devices forbid speculative or repeated accesses [2].

Architecture styles for application software do not, however, place these facts in the architecture:

- Layered and Clean Architecture [48], Ports and Adapters [17] and domain-driven design [22] constrain dependencies.
- Actors [4, 14, 33] and the Elm architecture [18] give state a single owner.
- Sagas [25] and the transactional outbox [63] keep state and messages consistent.

None of them says when an effect may leave a component, or what the component may claim after a timeout. Each feature author decides. In our benchmark (Section 8), straightforward implementations in four styles charged customers twice or reported charged orders as failed in 12–73% of runs under injected faults. A version hardened by hand, which writes the reconciliation procedure of our specification into the feature, was correct, but its correctness lived in the feature code.

This paper presents Ballarch, an architecture whose rules make that knowledge structural. Ballarch keeps established mechanisms (single owners, pure decision functions, commit then dispatch, adapters at the edge) and cites them as such. Its contributions are the following:

1. **Boundaries in time** (Section 3). Every interaction crosses two instants in order: the *commit horizon* and the *send horizon*. The instants are known from transaction processing, outboxes, transport protocols and durable execution (Section 9). We make them normative boundaries of a general architecture style, alongside the spatial ones: every implementation, whether a service, firmware or a circuit, must expose them as observable events, and the outcome rules are derived from them.

2. **Outcomes from phase and effect class** (Section 4). A runtime, not each feature, transport or adapter, chooses each request's single outcome from its durable phase and its port's effect class. It never reports NotDone without evidence and never re-sends a non-idempotent request. Transports and durable-execution engines already tell a request that never left from one that may have run [20, 23, 72, 73]. Ballarch derives the outcome from an effect class declared for each port, refuses retries of NonIdempotent ports by construction rather than by configuration, and gives the decider a typed three-valued completion that composes.

3. **A knowledge algebra with composition closure** (Section 5). Knowledge of a composite operation is a pair (*all*, *any*) over Kleene's strong logic, with absorbing compensations; it refines the heuristic-outcome merge of transaction managers. Because effects and compensations are encoded as read-once formulas, known properties of Kleene's logic make sound parts compose into a sound whole whose two coordinates are each exact when the calls are independent. A composite that counts what it may still do as unknown and replies only with settled knowledge is an honest constituent one level up, given honest adapters, absorbing compensations, and a call ledger kept for as long as the calls' effects persist.

4. **One rule set, checked by level-specific mechanisms** (Section 6). Fourteen rules are stated once, without reference to language, transport or storage. References for services (Kotlin),

firmware (C) and synchronous hardware (SystemVerilog) check the call-protocol rules by mechanisms suited to their level, some only in part; the Kotlin reference also checks the other rules that have a mechanism, including the algebra.

We evaluate the architecture with an Alloy model of the call protocol, exhaustive checks of the algebra, fault exploration at each level, and a benchmark comparing six architectural styles (Section 8). All artifacts are public [74].

2 Model

Balls and deciders. A *Ball* is a single-writer owner of state, addressed by a type and a key. Its type is defined by a *decider*, a set of pure functions:

- `decide(state, request, ctx)` returns *Accept*(state′, reply plan, outputs) or *Reject*(reason);
- `complete(state, completion, ctx)` and `observe(state, notice, ctx)` return a step (state′, outputs).

Outputs are *calls* to ports, *replies*, *notices* and *cancellations*. The context supplies time, identity and fresh call keys, so deciders read nothing ambient. The idea of a decider is not new [15, 18]; what is new is what the runtime does with its outputs.

Ports and effect classes. A *port* is a typed request/response contract owned by its destination. Its *effect class* states what repeating a request does there:

- Safe: no effect;
- Idempotent: the destination deduplicates by the call's identity;
- NonIdempotent: every delivery may take effect.

The names follow HTTP's safe and idempotent methods [23]. HTTP defines idempotence by the intended effect of repeated identical requests; Idempotent here requires that the destination deduplicate by the call's identity, as an idempotency key or token does [42, 67]. RPC interface languages already declare such classes per operation to decide client retries [61, 67]; Ballarch also derives each call's outcome from its class (Section 4). In hardware the same distinction appears as Normal versus Device memory [2], and as prefetchable PCIe memory ranges, whose reads have no side effects [59].

Obligations and outcomes. A committed call is an *obligation*: the runtime must deliver exactly one completion for it to the caller's `complete`, by a finite deadline. The completion's outcome is one of the following:

- Done(value): the destination processed the request; the value may describe a business refusal, such as a declined charge;
- NotDone(reason): there is evidence that the request had no effect;
- Unknown(reason): the request may or may not have taken effect.

Profiles. Under the *Transient* profile nothing survives a crash. Under the *Durable* profile the committed state survives, together with the call ledger and its phases, recently answered requests and unacknowledged notices.

3 Boundaries in Time

Spatial boundaries constrain which code may refer to which. Ballarch adds two boundaries that constrain *when* effects and knowledge may cross (Figure 1):

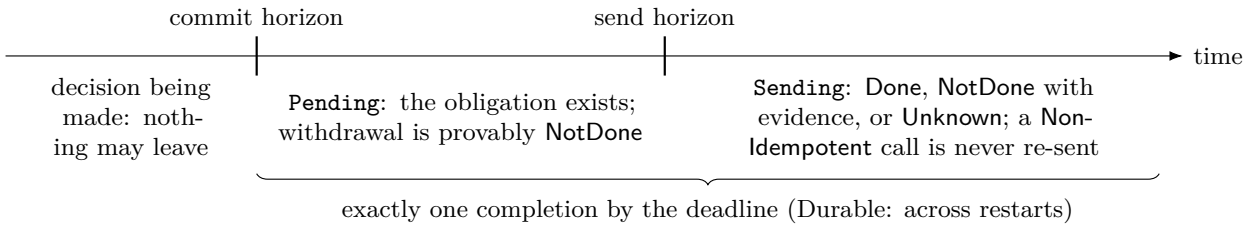


Figure 1: The two boundaries in time that every call crosses. Before the commit horizon there is nothing to know; between the horizons non-occurrence can still be proved; after the send horizon only evidence about every attempt, or later reconciliation, establishes what happened to an effectful call.

Commit horizon. The atomic commit of the decision that requests a call. Before it, the call is part of a decision that can still vanish, and nothing may leave (commit before dispatch). After it, the call is an obligation.

Send horizon. The first moment the request may have left the runtime. Before it, the call can be withdrawn truthfully. After it, NotDone requires evidence covering every attempt (a refusal, or an adapter's evidence of non-application), except for a Safe call, which can have had no effect.

An implementation must make each horizon an event that its checks can observe:

- in the Kotlin runtime: a store write, then a write-ahead `Sending` record written before the transport is invoked;
- in firmware: the CRC-checked NVM record, then the write-ahead record before the UART transmission;
- in a synchronous circuit: the clock edge, then the bus handshake.

The mechanisms behind both horizons are established:

- the output-commit rule of rollback recovery [21, 68];
- write-ahead logging [51];
- operating systems that hold external output until it no longer depends on uncommitted state [54, 55];
- in-order retirement in processors [66].

Real-time architectures specify component interfaces in the temporal domain [40], but for timing predictability rather than for knowledge about effects. The instants themselves are known in practice. Transaction processing defers real actions, which cannot be undone, until commit [27], and the transactional outbox commits messages together with state [63]. Service-design guidance commits each local atomic phase before a foreign state mutation, and counts the request as committed once the first such mutation is made [42]. Durable-execution engines record when an activity is scheduled and when it is dispatched to a worker [70]. What we add is to make both instants normative boundaries of a general architecture style, which every implementation must expose as observable events, and to derive from them outcome rules that keep Unknown apart from failure (Section 4). The service-design guidance above, for example, marks a timed-out non-idempotent call as failed. Section 9 compares the outcome rules with the retry rules of transports and durable execution.

4 Outcomes from Phase and Effect Class

Table 1 gives the classification that a conforming runtime applies. Two rules carry its weight:

Situation	Safe	Idempotent	NonIdempotent
Destination answered	Done	Done	Done
Deadline or cancel before the send horizon	NotDone	NotDone	NotDone
Evidence of non-application, one attempt	NotDone	NotDone	NotDone
Same, after more than one attempt	NotDone	retry or Unknown	–
No answer, attempts and time left	retry	retry	Unknown
No answer, no attempts or time left	NotDone	Unknown	Unknown
Durable restart after the send horizon	re-send or NotDone	re-send or Unknown	Unknown
Durable restart before the send horizon	send	send	send

Table 1: The runtime's classification of a call's single outcome. A Safe call that gets no answer ends NotDone(NoAnswer), because it can have had no effect. Evidence that one attempt was not applied says nothing about earlier attempts, so after several attempts of an Idempotent call it does not justify NotDone; a NonIdempotent call has only one attempt. A report that arrives after the call resolved is dropped.

R7 (honest outcomes). NotDone is produced only with evidence: the call was never sent, the destination refused it, an adapter reported evidence covering every attempt that may have been delivered, or the port is Safe and so has no effect.

R8 (no blind re-send). A NonIdempotent call is sent at most once, including after a restart. Before each send, the runtime durably records that the call may have been sent.

Retries belong to the runtime and are allowed only for Safe and Idempotent ports. Adapters report each attempt once and never retry. Table 2 lists all fourteen rules with the mechanisms that check them.

The decider sees one completion per call and must handle all three outcomes. The type system makes Unknown a case, not an exception that each call site maps differently:

```
when (val outcome = completion.outcomeOf(Mailer)!!) {
    is Outcome.Done -> step(ReminderState.Sent)
    is Outcome.NotDone -> step(ReminderState.NotSent(outcome.reason.toString()))
    is Outcome.Unknown -> step(ReminderState.MaybeSent) // never re-send blindly
}
```

5 Knowledge of Composite Operations

A Ball often performs one operation through several calls: a flight and a hotel, a reservation and a charge. Workflows usually report such operations as succeeded or failed. After a partial or uncertain effect, both labels can be false.

Definition 1 (Truth values)**.** Let $\mathbb{T} = \{1, 0, \bot\}$ with Kleene's strong connectives [39] and the information order $\bot \sqsubseteq 1$, $\bot \sqsubseteq 0$. A value t *admits* a Boolean b if $t = \bot$ or t is the truth value of b (1 for true, 0 for false). A pair of values admits a pair of Booleans if each coordinate admits the corresponding Boolean.

Definition 2 (Constituents and knowledge)**.** The *constituents* of an operation are the calls it needs, requested or not. A claim about the operation is about its *final* effects, those that hold once none of its requests can still take effect. At any moment, the *effect* $f_i \in \mathbb{T}$ of a constituent is:

Rule	Binds	Statement (abridged)	Checked in this work by	Levels
R1 Single owner	author	one owner per fact; immutable values	bytecode rules; C object-file check; Yosys single-driver check	K C R
R2 Pure decider	author	no I/O, clocks, randomness, globals	bytecode rule; C undefined-symbol check; Yosys latch check	K C R
R3 Serial processing	runtime	no concurrent or nested steps per instance	runtime guard; trace oracle	K
R4 Atomic commit	runtime	state and outputs in one atomic write	store design; Alloy P6; power-loss exploration	K C
R5 Commit before dispatch	runtime	nothing leaves before its commit	Alloy P1; oracles; power-loss exploration; Yosys path check	K C R
R6 One completion per call	runtime	exactly one by the deadline	Alloy P2, L1; oracles; C exploration; RTL bench	K C R
R7 Honest outcomes	runtime, adapter	NotDone only with evidence	Alloy P3, P5; oracles with ground truth; C exploration; RTL bench (no false NotDone)	K C R*
R8 No blind re-send	runtime, adapter	NonIdempotent sent at most once; write-ahead	Alloy P4; oracles; seeded bugs; C exploration; RTL bench	K C R
R9 One reply per request	runtime, author	at most one reply; overdue replies reported	preflight; oracle	K
R10 Explicit composition	author	interaction only via calls, replies, notices	kernel API; module-boundary check	K
R11 Bounds and preflight	runtime	refuse before accepting; reject faulty decisions	preflight	K
R12 Read-only queries	author	reads do not change state	types	K
R13 Honest profile claims	runtime	no claim beyond the profile	review	–
R14 Honest composition	author	claims about composites follow the algebra	`Knowledge` type; exploration against ground truth	K

Table 2: The rules of Ballarch Core 2.0 and the mechanisms that check them. Levels: the references that check the rule mechanically (K Kotlin service, C firmware, R RTL). R*: in part; the RTL bench has no fate in which an accepted command is neither applied nor answered, so it checks that NotDone is never reported for an applied command, but not that an unanswered command ends Unknown.

- 1 if it completed Done with an answer that, by the port's contract, means that the requested effect happened;
- 0 if it completed NotDone, if it completed Done with an answer meaning that the effect did not happen (a recorded refusal), or if the Ball can no longer request it for this operation;
- $\perp$ if it completed Unknown, if it completed Done with an answer that leaves the effect open, if it was requested and its completion has not arrived, whether or not it was sent, or if the Ball has not requested it but still may.

If the Ball may request the same effect again after a NotDone, its attempts, including those it may still make, form one constituent, whose effect f_i is the disjunction of theirs, each defined as above. A compensation is *absorbing* if its effect means that no effect of the original request exists or can still take effect. For example, the destination fences the request's reference, as the anti-suspension fence of Try-Confirm/Cancel implementations does [1], or the original request's *delivery horizon*, the age after which the destination no longer applies it, has passed. In Pardon and Pautasso's design,

a participant cancels a reservation that is not confirmed within its timeout and refuses a later confirmation [57]; that timeout is the confirmation's delivery horizon. Let $c_{i1}, \ldots, c_{im}$ be the effects of a constituent's compensation attempts, each defined like f_i, and let u_i be their disjunction, joined with $\bot$ while the Ball may still compensate the constituent. If the compensations are absorbing, the constituent's *net effect* is $n_i = f_i \wedge \neg u_i$. The *knowledge* of an operation with constituents $1..k$ is

$$K = (all, any) = \Big(\bigwedge_i n_i,\ \bigvee_i n_i \Big), \qquad K_\emptyset = (1, 0).$$

Counting every call that the Ball may still make as $\bot$ makes K a claim about final effects at every moment, not only at the end. A report of the completions received so far is a different claim, and must not be offered as the operation's outcome. An operation is *settled* when none of its calls is open and none can still be made; from then on, only Unknown completions and open answers leave a constituent at $\bot$.

Composition $(a_1, s_1) + (a_2, s_2) = (a_1 \wedge a_2,\ s_1 \vee s_2)$ is a commutative monoid with identity $K_\emptyset$, the knowledge of an operation that requested no effect. A constituent contributes the summand (n_i, n_i). A call to a composite port contributes a pair instead: the composite's settled knowledge if the call completed Done with it, $(0, 0)$ if it completed NotDone, and $(\bot, \bot)$ if it completed Unknown, was answered before the composite settled, is still open or may still be made. Seven values can occur:

- Done $(1, 1)$ and NotDone $(0, 0)$;
- Partial $(0, 1)$: some effects happened and some did not;
- NotAllDone $(0, \bot)$ and SomeDone $(\bot, 1)$;
- Unknown $(\bot, \bot)$;
- NoEffects $(1, 0)$.

Seen as one call, a composite reports $\pi(K)$: Done if $K = (1, 1)$, NotDone if $any = 0$ (this includes NoEffects), and Unknown otherwise.

Both propositions below rest on standard facts. Kleene's strong connectives are monotone in $\sqsubseteq$ [39]. On formulas in which every variable occurs once, they agree with supervaluation, in van Fraassen's sense [76]: the value is definite exactly when all ways of resolving the unknowns agree. This agreement is folklore; the proof of Proposition 2 includes it. Hardware engineers know the second fact from ternary (0, 1, X) simulation, which goes back to hazard detection in switching circuits [19]. Read $\bot$ as Verilog's `x`: the proof of Proposition 2 also shows why X-propagation is exact through logic without reconvergent fanout, and why it can be pessimistic with it. What is ours is the encoding of effects and absorbing compensations as such formulas, and the closure property that follows from it.

Proposition 1 (Soundness)**.** *Let b_i be the true final net effect of constituent i. If each n_i admits b_i, then* all *admits $\bigwedge_i b_i$ and* any *admits $\bigvee_i b_i$. The same holds if some summands are pairs (a_i, s_i), each admitting the conjunction and the disjunction of the true net effects of its own constituents.*

Proof. Kleene's strong connectives are monotone in $\sqsubseteq$ and agree with the Boolean ones on definite values. Replacing each $\bot$ among the n_i by the corresponding b_i refines every n_i to b_i, so $\bigwedge_i n_i \sqsubseteq \bigwedge_i b_i$, which is the claim for *all*; the claim for *any* follows in the same way. For pairs, each coordinate of $+$ is monotone, and by associativity the conjunction of the inner conjunctions is the conjunction over all inner constituents; the same holds for disjunctions. □

Proposition 2 (Coordinate-wise exactness)**.** *Let the effect of each call, including each repeated attempt and each compensation attempt, be the most informative value that admits every possibility for that call alone. Assume that the calls are* independent*: the combinations of true effects consistent*

with the outcomes are exactly the product of the per-call sets of possibilities. Then each coordinate of K is the most informative value that admits every such combination.

Proof. For a formula φ over the calls' effects, let $S(\varphi) \subseteq \{0,1\}$ be the set of its Boolean values over the consistent combinations. The most informative value admitting all of them, the supervaluation, is $\mathsf{1}$, $\mathsf{0}$ or $\bot$ as $S(\varphi)$ is $\{1\}$, $\{0\}$ or $\{0,1\}$. We show by induction that, on a formula in which each variable occurs at most once, Kleene evaluation equals supervaluation. For a variable, this is the hypothesis. A call that the Ball may still make enters as a variable; since the Ball's later calls often depend on earlier outcomes, independence then usually holds only once the operation is settled. For a constant it is immediate, which covers $K_\emptyset$. For $\neg\varphi$, $S(\neg\varphi) = \{\neg x : x \in S(\varphi)\}$. For $\varphi_1 \circ \varphi_2$ with $\circ \in \{\wedge, \vee\}$, the two sides share no variable, so by independence $S(\varphi_1 \circ \varphi_2) = \{x \circ y : x \in S(\varphi_1),\ y \in S(\varphi_2)\}$, and the nine cases are Kleene's tables. *all* and *any*, with each $n_i = f_i \wedge \neg(c_{i1} \vee \cdots \vee c_{im})$ substituted, and each f_i of repeated attempts written as the disjunction of their effects, are such formulas, because every attempt is a distinct call. Both hypotheses are needed. With a repeated variable, $x \vee \neg x$ evaluates to $\bot$ although its supervaluation is $\mathsf{1}$; in a circuit, this is reconvergent fanout. With correlated calls, the product step fails: if a Ball books a hotel only after its flight ends NotDone, then while the flight is open *all* evaluates to $\bot$, although no consistent combination has both effects, so its supervaluation is $\mathsf{0}$. □

Corollary 1 (Composition closure)**.** *Assume that:*

(i) adapters and destinations report honestly (R7, and Done answers follow the port's contract);
(ii) every compensation counted in a net effect is absorbing;
(iii) the Ball keeps its state and call ledger for as long as the effects of its calls persist: under the Durable profile, or under the Transient profile only if a crash also clears those effects;
(iv) the Ball derives every claim about a composite operation from K (rule R14), and answers the request that started the operation only once the operation is settled (a deferred reply*), with K itself or with $\pi(K)$.*

Then:

(a) at every moment, K admits the conjunction and the disjunction of the constituents' true final net effects;
(b) the reply admits them as well;
(c) one level up, the summand that the call to the composite contributes to the caller's knowledge, as described after Definition 2, admits them too.

By induction on the depth of nesting, every claim derived this way, at every level, admits the truth.

Proof. (a) By (iii), every call made for the operation whose effect can still exist is recorded in the Ball's state, so it is counted in some f_i or u_i. By (i), the value of each such call admits its true effect at every moment: honest reporting covers Done and NotDone, a call that can no longer be made never takes effect, and every other case is $\bot$. By monotonicity, a disjunction of repeated attempts admits whether one of them takes effect, so each f_i admits whether its constituent's request takes effect. While the Ball may still compensate, the disjunction with $\bot$ keeps u_i at $\bot$ unless a compensation already took effect, so u_i admits whether some compensation of constituent i takes effect. Because the compensations are absorbing (ii), the true net effect is $b_i = \beta_i \wedge \neg\gamma_i$, where β_i says whether the constituent's request takes effect and γ_i whether one of its compensations does. By monotonicity, n_i admits b_i, and Proposition 1 gives (a). (b) By (iv), a reply with K is (a) at the moment of settlement. For $\pi(K)$, let the composite's true effect as one call be any Boolean B with $\bigwedge_i b_i \le B \le \bigvee_i b_i$ when $k \ge 1$, such as 'every effect happened' or 'some effect happened', and $B = 0$ when $k = 0$. If $\pi(K) = \mathsf{Done}$, then $all = \mathsf{1}$, so $\bigwedge_i b_i = 1$ and $B = 1$. If $\pi(K) = \mathsf{NotDone}$, then $any = \mathsf{0}$, so $\bigvee_i b_i = 0$ and $B = 0$. Otherwise $\pi(K) = \mathsf{Unknown}$, which admits B. (c) A Done

completion carries the reply, which (b) covers, read as a pair or as one effect. A NotDone completion of a call to a Ball means that the request was refused or never delivered, so the composite made no call for it, which $(0, 0)$ admits; $(\bot, \bot)$ admits anything. The summand therefore satisfies the hypothesis of Proposition 1 one level up, and induction on the depth of nesting completes the proof. □

The counting rule and the deferred reply both matter. A reply computed this way before settlement is sound, but it says little, because the calls still to come are $\bot$. A reply that counted them as 0 would not be sound: a trip that answers at once that its hotel is not booked is contradicted when the hotel is booked a moment later. A deferred reply takes time, so a caller whose deadline is shorter than the composite's settlement ends its call Unknown, which is honest. Hypothesis (iii) of Corollary 1 matters as well. Under the Transient profile, a trip that crashes after its flight request was sent restarts with an empty ledger. If its caller sends the request again, the new trip books again and can settle with NotDone, for example after a sold-out answer, while the first seat is held. The projection π is sound but not exact: Partial, NotAllDone and SomeDone all become Unknown. A composite that can end partially should therefore reply with K. The caller then holds, by associativity, the knowledge of the flattened operation, to which Proposition 2 applies when all leaf calls are independent.

Transaction managers already merge and nest outcomes, more coarsely. Participants report heuristic outcomes [77]. The Narayana transaction manager, for example, merges one participant's heuristic commit and another's heuristic rollback into *mixed*, which absorbs later outcomes, and its interposed subordinate coordinator reports the merged outcome to its superior as one participant [53]. *Hazard* does not record which way the known participants went, so it covers NotAllDone, SomeDone and Unknown: merging hazard with a rollback stays hazard, whereas SomeDone + NotDone = Partial. In that implementation the merge also depends on arrival order: commit, hazard, rollback gives hazard, while commit, rollback, hazard gives mixed [53]. Here + is commutative, and both orders give Partial. Heuristic outcomes also signal a failure of atomicity, whereas the algebra describes operations that are not atomic, for which Partial is an ordinary result.

The kernel's `Knowledge` type implements the algebra. An exhaustive test checks three things:

- the connectives against supervaluation, which is the inductive step of the proof of Proposition 2;
- the monoid laws on all seven values;
- soundness and exactness for up to four constituents, each with at most one compensation: 69,905 worlds.

The propositions hold for any number of constituents and compensation attempts; the test checks small cases directly.

Exactness is relative to the constituents' individual outcomes. If $\mathbf{K}e$ reads "it is known that e", conjunction matches epistemic logic, since $\mathbf{K}(a \wedge b) \Leftrightarrow \mathbf{K}a \wedge \mathbf{K}b$ [30], but disjunction does not. Evidence that correlates constituents, such as a log showing that one of two calls landed without saying which, can establish more than the algebra derives. The algebra under-approximates such knowledge and never over-approximates it.

Example. A trip *Flow*, a Ball whose state is a multi-step workflow, books a flight, which is NonIdempotent, and a hotel, which is Idempotent. It compensates with absorbing voids when it cannot hold both, and answers the booking request only once settled, with its final K. Under the Durable profile, the exploration ran 1,339 runs within a bound of twelve choices per run, covering every combination of five fates per booking attempt, three per compensation attempt, and one crash before any of the first ten scheduler tasks. In every run that accepted the booking, its settled knowledge admitted the providers' real bookings. The resulting knowledge values were Done 2,

Concept	Service (Kotlin)	Firmware (C11)	Hardware (SystemVerilog)
Owner (R1)	runtime-held state	one NVM record	registers with one writing process
Decider (R2)	pure functions	`step()` in its own unit	one latch-free `always_comb`
Commit horizon	one store write	CRC-checked slot write	the clock edge
Send horizon	write-ahead `Sending`	write-ahead record	`valid && ready` (AXI: VALID [3])
Unknown	timeout or restart while `Sending`	power loss while `Sending`	accepted transaction without response
Profile	Transient or Durable	Durable	Transient

Table 3: The same concepts at three levels.

NotDone 282, NotAllDone 620, Partial 64, SomeDone 92, and Unknown 276. On the same runs, a two-valued status ("booked" if both bookings answered booked, otherwise "failed") was wrong in 254 runs. An absorbing compensation that completes Done turns an Unknown booking into a definite NotDone, since $\bot \wedge \neg 1 = 0$; a compensation whose answer is lost leaves the booking unknown. Compensation narrows ambiguity; it does not remove it.

6 One Rule Set at Three Levels

The rules mention no language, transport or storage. Table 3 maps the concepts onto the three references; each reference also carries seeded bugs that its checks must catch. The references do not check the same rules: the last column of Table 2 lists the levels at which each rule is checked.

Services. The Kotlin reference consists of the following parts:

- a dependency-free kernel of about 260 lines of code, including the knowledge algebra;
- a runtime of about 1,000 lines that implements both profiles, with an in-memory store for simulation and a file store that commits by atomic rename;
- a test kit whose oracles check R3–R9 on execution traces against simulated ground truth, under exhaustive or seeded exploration of network fates and crash points in the style of stateless model checking [26];
- architecture rules that check R1, R2 and R10 on bytecode.

Firmware. The C reference commits state and the call ledger as one CRC-checked record in two NVM slots. It records each send ahead of transmission, so a power loss after the write-ahead record turns a dose command into Unknown rather than a second dose. Its example is an infusion-pump controller whose pump has no idempotency key. A bounded exhaustive exploration covers about 220,000 runs over link fates and power-loss points, including torn writes. No violation of R5–R8 occurred, nor any overdose or mismatch between the controller's belief and the pump. A check of the decider's object file rejects undefined symbols (R2) and writable data (R1). It is necessary but not sufficient, since memory-mapped I/O and inline assembly leave no symbol. Intermittent computing meets the same hazard on energy-harvesting devices (Section 9).

Hardware. The RTL reference is a command issuer:

- its registers have one writer;
- its next-state logic is one latch-free `always_comb` block;

- its outputs depend on registered state only.

It issues READ (Safe), CONFIG (Idempotent) and FIRE (NonIdempotent) transactions on a bus with id-tagged responses. Yosys checks single drivers, the absence of latches, and the absence of any combinational path from an input to a dispatched output. A fault-injecting bench runs 75 scenarios against a device model. In that model, every accepted command that gets no answer in time has been applied, so the bench checks that NotDone is never reported for an applied command, but not that an unanswered command ends Unknown rather than Done. The first draft of the block reported NotDone when a CONFIG retry was withdrawn before acceptance, although the first attempt had already been applied; the bench caught it. Hardware engineers apply such rules by experience. Stating them as rules lets a tool check them.

Closest work. Bluespec's guarded atomic actions and the Bluespec Codesign Language compile one rule-based semantics to both hardware and software [36, 56]. Lingua Franca reactors and synchronous languages give deterministic semantics across targets [7, 29, 44]. These unify how computation executes. Ballarch unifies what is known about external effects, and checks its call protocol at each of the three levels.

7 Formal Model of the Call Protocol

An Alloy 6 model [35] covers the protocol of Table 1:

- commit, write-ahead, send, retry, deadline, crash and recovery;
- a lossy, duplicating network;
- the three effect classes.

Within bounds of three calls and traces of up to 14 steps, it has no counterexample to six safety properties: no send before commit, stable outcomes, no false NotDone, at most one send of a NonIdempotent call, Done only after processing, and agreement of state and call ledger. Under fairness, with two calls and up to 12 steps, every awaited call gets an outcome. Each of six seeded runtime bugs yields a counterexample, for example sending before the write-ahead record or re-sending after a crash. These are bounded results, not proofs for all sizes.

8 Evaluation

We ask three questions:

RQ1. Under provider faults and crashes, do implementations that follow each style's canonical guidance meet the correctness requirements?

RQ2. What does Ballarch cost and gain structurally, compared with a design hardened by hand to the same correctness?

RQ3. Do the oracles detect seeded bugs in the correct implementations?

Scenario. Three orders compete for two units of stock. Each order reserves a unit, charges 100 through a provider *without idempotency keys*, and completes; on failure it releases the unit. The provider's status endpoint is Safe. A database with atomic transactions survives crashes. The requirements are:

- no double charge;
- no order reported failed while charged, and none reported completed while not charged;

Architecture	Stress	Stress + crash	Mild	Mild + crash
Layered (Clean Architecture)	73.4	67.3	11.7	57.9
Layered, hardened by hand	0.0	0.0	0.0	0.0
Actors with durable state	73.4	68.2	11.7	35.9
MVU (Elm architecture)	73.4	50.7	11.7	48.5
Hexagonal + outbox + saga	73.4	72.5	11.7	24.9
Ballarch	0.0	0.0	0.0	0.0

Table 4: RQ1: percentage of runs (of 2,000 per cell) that violated at least one requirement. Wilson 95% intervals have a half-width of at most 2.2 points; for 0.0 the upper bound is 0.2. Without crashes, the four idiomatic designs behave identically: each maps a timeout or error to failure, with one attempt per order.

- every accepted order ends in a terminal state;
- no stock leak, oversell or charge for an unaccepted order.

A NeedsAttention terminal state is allowed and is reported as a cost.

Fault model. Each provider request draws one fate: answered in time, answered late, lost, processed with the answer lost, processed then reset, or refused before sending.

- *Stress* draws all six fates with equal probability.
- *Mild* answers 90% of requests in time.

In crash cells, the process crashes once, before a scheduler task chosen uniformly within a typical run of that architecture, and restarts 50 ms later. The simulation is deterministic; rerunning it reproduces the published data byte for byte.

Implementations. The same author wrote all six against this specification. Code that stands in for a framework is excluded from the structural metrics:

Layered Clean Architecture [48]. One attempt with a timeout; any exception fails the order. On restart it resumes unfinished orders.

Hardened The same layers, plus a hand-written write-ahead status, retries only after a refused connection, and reconciliation by status query after the provider's delivery horizon: the reconciliation procedure of the specification's guidance, written into the use case.

Actors Durable state and `ask` with a timeout. On restart, actors are re-driven.

MVU The Elm architecture [18] on the server. The model is persisted before commands run, and `init` re-issues them.

Saga Hexagonal architecture with a transactional outbox, an orchestrated saga and idempotent consumers [63].

Ballarch Inventory, payment and checkout Balls on the reference runtime.

RQ1. Table 4 shows the results. The violations of the idiomatic designs come from two sources:

- a timeout or reset after the charge took effect, reported as a failure (false failures);
- recovery paths that re-send a charge that may already have happened (double charges): resumed layered processing, re-driven actors, re-issued MVU commands, and outbox redelivery to a consumer whose provider call lies outside its transaction.

In the Mild-with-crash setting, for example, the layered design double-charged in 1,110 runs and the saga in 338. Actors also leaked stock when a crash lost a fire-and-forget release. Ballarch and

Metric	Layered	Hardened	Actors	MVU	Saga	Ballarch
SLOC	166	216	163	96	184	279
Classes	12	20	15	10	18	35
CBO, mean (max)	4.00 (9)	4.70 (16)	3.47 (14)	4.00 (9)	4.11 (13)	3.94 (18)
LCOM4, share = 1	71%	71%	67%	50%	67%	8%
CAMC, mean	0.52	0.52	0.39	0.42	0.57	0.33
Distance D, mean	0.34	0.35	0.42	0.39	0.47	0.44
Propagation cost	28.5%	20.0%	21.8%	31.0%	20.7%	13.1%
McCabe CCN, mean (max)	2.05 (8)	1.81 (7)	1.60 (4)	1.75 (7)	2.46 (13)	2.19 (8)
McCabe CCN, total	41	49	32	21	32	68
Framework types	5	6	5	3	6	45
`catch` clauses	3	5	0	0	2	0
Code in pure classes	26%	27%	26%	59%	44%	88%

Table 5: RQ2: structure of the feature code. Metrics: CBO [16], LCOM4 [34], CAMC [5], distance from the main sequence [47], propagation cost [45], and McCabe complexity [49] via lizard [79]. Structural metrics come from bytecode; compiler-generated classes are merged into their declaring class.

the hardened design had no violation in any of the 16,000 runs between them. The cost of their honesty is NeedsAttention: 0.27–0.31 per order under Stress, and at most 0.001 under Mild.

RQ2. Table 5 compares the structures. Against the hardened design, which is equally correct, Ballarch has these costs:

- 29% more source lines and 28% more bytecode;
- 45 framework types against 6;
- more, smaller classes, with higher cyclomatic complexity (2.19 against 1.81 per function, 68 against 49 in total);
- a larger mean distance from the main sequence (0.44 against 0.35), mostly from its shared `api` package of concrete message types ($D = 0.79$), and the most coupled single class of the six designs (CBO 18 against 16);
- lower cohesion.

LCOM4 assumes that cohesive methods share fields. It counts stateless deciders over a state type as uncohesive, which affects MVU's `update` as well. The type-based CAMC metric does not reverse the result. In exchange, Ballarch has a lower mean CBO than the hardened design (3.94 against 4.70) and the lowest propagation cost of the six designs. It also places 88% of its code in classes with no mutable field and no reference to concurrency, I/O, clocks, the database or the provider, against 27% for the hardened design. Its feature code has no exception handler. The hardened design reaches the same correctness through five `catch` clauses, a write-ahead status and reconciliation, all written into the use case.

RQ3. We injected two plausible bugs from outside the correct implementations: a Ballarch payment Ball that treats Unknown as not charged, and a hardened gateway that reports an uncertain charge as not sent. Under Stress with a crash they were detected in 79.4% and 46.4% of runs respectively, and under Mild with a crash in 47.3% and 2.3%.

9 Related Work

Architecture styles. Architecture research describes styles as components, connectors and constraints [60, 65], and modules by the decisions they hide [58].

- Layered, Clean and Hexagonal architectures and domain-driven design [17, 22, 48] constrain dependencies.
- Actor systems give each piece of state one owner [4, 14, 33].
- The Elm architecture and the Decider pattern separate pure transitions from effects [15, 18].
- Durable-execution systems persist the progress of orchestrations and replay them deterministically [12].

Ballarch keeps the structural constraints of these styles. To them it adds the temporal boundaries, the outcome rules and the algebra. Meng and Jackson's structural pattern for legible software is a recent example of a new structural pattern evaluated on an application [50].

Ambiguous outcomes and idempotence. Several lines of work recognise ambiguous outcomes:

- the third outcome of RPC, a failed call that may or may not have been executed [8];
- Jepsen's indeterminate operations [37, 38];
- the heuristic outcomes of XA [77], including mixed and hazard;
- Helland's analysis of activities between partners [31, 32].

X-ability defines exactly-once correctness for services with side effects and is composable [24]. Ramalingam and Vaswani make workflows idempotent and therefore safe to retry [62]. These approaches hide or remove the third outcome. Recent studies of LLM agents report duplicated and hidden side effects at tool boundaries [43, 75], a new setting for the same problem.

Retry rules in transports and runtimes. Transport protocols already apply send-horizon reasoning to single requests. HTTP says that a client should not automatically retry a request with a non-idempotent method unless it knows that the request is in fact idempotent, or can detect that the original request was never applied [23, Sec. 9.2.2]. HTTP/2 supplies such evidence: a GOAWAY frame's last stream identifier and the REFUSED_STREAM error code guarantee that a request was not processed, and a client may then retry it even if it is not idempotent [73, Sec. 8.7]. HTTP/3 defines a rejected request in the same way [9, Sec. 4.1.1]. gRPC retries transparently only while a call has not left the client or has not reached the server's application logic [20]; it calls an RPC *committed* when it stops making new attempts, an instant unrelated to our commit horizon. Service-design guidance commits local state before each foreign mutation, retries only mutations that are idempotent or carry a key, and marks a timed-out non-idempotent one as failed [42]. Durable-execution engines run activities under a retry policy, with unlimited attempts by default in Temporal; an attempt that times out after dispatch is retried if the policy allows it [71, 72]. Temporal records which timeout ended an attempt, so a workflow can tell an attempt that never reached a worker (Schedule-To-Start) from one that may have run (Start-To-Close) [70, 72]. Mapping that to an outcome, and not retrying a non-idempotent activity, are left to workflow code and retry-policy configuration.

These rules govern the retries of one protocol, client or service, and report the rest as errors or leave them to configuration. Ballarch makes them one contract of the architecture. The runtime derives each outcome, and whether it may retry, from the effect class declared for the port; it never retries a NonIdempotent call (R8). The outcome is a typed value delivered once to the owner's committed state, Unknown is kept apart from failure, under the Durable profile the phase survives crashes through the write-ahead record, the result composes through the algebra of Section 5, and the same rule is checked in firmware and RTL. Ballarch is also more conservative than HTTP/2:

evidence of non-application ends the call as NotDone, and the decider, not the transport, decides whether to send it again.

Transactions and fault-tolerant runtimes. Transaction processing defers real actions until commit [27], and a participant in atomic commitment is uncertain of the outcome between its vote and the decision [6]; the two horizons carry these instants from transactions to every call of an application. Reliable State Machines are actor-like state machines whose state survives failures, with a proof of failure transparency [52]; they are the closest runtime model to Balls. Such systems aim at exactly-once processing inside their boundary. Ballarch addresses destinations outside any such boundary, which may offer no key: it turns the remaining uncertainty into a typed outcome and composes it. The end-to-end argument [64] places duplicate suppression and delivery assurance at the endpoints. The Idempotent class puts deduplication at the destination, and the outcome rules let the caller state only what it can know.

Compensation and three-valued reasoning. Sagas [25] and compensation calculi [10, 13] give semantics to long-running transactions, but assume that outcomes are known. Try-Confirm/ Cancel [57] does not: its implementations fence a request whose cancel arrived first ("empty rollback" and "anti-suspension" [1]), which is our absorbing compensation. Transaction managers merge heuristic outcomes and pass them up through interposed coordinators [53]. They thus say what a composite may claim while some outcomes stay unknown, but coarsely, depending on arrival order, and only for atomic transactions; Section 5 shows how the (*all*, *any*) algebra refines that merge. None of these approaches counts calls still to be made, or compensations, in such a claim. Three-valued abstractions and may/must approximations [11, 41] use the monotonicity argument of our soundness proposition, and ternary circuit simulation [19] uses the same connectives; its pessimism under reconvergent fanout is why exactness needs read-once formulas. Knowledge-based analyses of commitment [28, 30] reason about what processes know. The algebra here is truth-functional by design, so that a runtime can compute it.

Output commit and hardware. The output-commit problem [21, 68], write-ahead logging [51], external synchrony [54, 55] and precise interrupts [66] are the mechanisms behind the horizons. Intermittent computing meets the firmware reference's hazard on energy-harvesting devices: repeating I/O after a power failure causes idempotence bugs [69], peripherals need just-in-time checkpoints [46], and EaseIO lets programmers annotate each I/O call with its re-execution semantics [78]. Those annotations are the closest firmware analogue of effect classes. These systems aim to make re-execution correct; Ballarch instead never re-executes a NonIdempotent command whose fate is unknown, reports Unknown, and applies the same rule at the other levels. The time-triggered architecture [40] makes time part of interfaces. Bluespec, BCL, Lingua Franca and synchronous languages [7, 29, 36, 44, 56] unify execution semantics across hardware and software.

10 Limitations and Threats to Validity

- **Construct and internal validity.** One author wrote all implementations and designed the architecture. Each idiomatic baseline is one reading of its style's guidance, and the idiomatic baselines deliberately add no fault handling beyond it. The hardened baseline hand-implements the write-ahead status, evidence-based retry and reconciliation procedure of the Ballarch specification. It shows that correctness is reachable without the runtime, not how an independent expert

would harden the code. Independent implementations, written before seeing the results, would strengthen the comparison. The benchmark was not pre-registered.

- **External validity.**
  - The scenario is small and uses a simulated provider without idempotency keys; with idempotency keys, simple retries would become correct.
  - The fault distributions are assumptions.
  - The firmware and RTL references are single examples checked in simulation. They check the call-protocol rules, not every rule (Table 2).
  - No reference composes levels, for instance firmware that drives the RTL block as a Ball behind a port, so closure across levels is proved for the algebra but not exercised.
  - The Kotlin runtime's file store is single-node and has been tested against killed processes, not power loss; it is not a production database.

- **Metrics.** LCOM4 and CAMC were designed for object-oriented code with shared fields. Propagation cost and CBO depend on how finely code is divided into classes. Change impact, developer effort and runtime overhead were not measured in this study. An earlier prototype comparison, whose data is not published, found that splitting each feature into an API module and an implementation module made cross-feature changes touch more files than single-module baselines in 4 of 5 change requests, and that a local request took a median of 39 µs more than a direct call, almost all of it thread hand-off.

- **Formal results.** The Alloy results are bounded. The algebra's propositions and the closure corollary are proved, and the propositions are also checked exhaustively for small sizes. The closure corollary assumes honest adapters, absorbing compensations, a call ledger kept for as long as the calls' effects persist (the Durable profile, or crashes that also clear those effects), and deferred replies; the first two are reviewed, not verified.

11 Conclusion

Ballarch places two instants that mechanisms in recovery, operating systems and processors already respect into the architecture itself, as boundaries in time. A runtime derives each outcome from them, and a small algebra carries honesty from parts to wholes that reply with settled knowledge. The call-protocol rules can be checked by level-specific mechanisms from register-transfer logic to services. In our benchmark, correctness under faults required either Ballarch or the same reconciliation procedure written by hand into the feature. Ballarch moved that knowledge from the feature into the architecture, at the price of more code and more concepts. Whether that trade pays off for teams, and for coding agents that write such features, is the next question to measure.

Artifacts. The specification, the Kotlin, C and SystemVerilog references, the Alloy model, the benchmark and its raw data are available in the Ballarch repository, `https://github.com/4wl2d/Pokeball`, which still carries the project's earlier name (version 2.0.0-rc.1). The file `docs/problems.md` there maps twelve recurring failures of current architectures to the mechanism, rule and evidence for each, and lists what Ballarch does not solve. Text is licensed under CC BY 4.0; the code has no software license yet.

Use of generative AI. The author used an AI assistant (Anthropic's Claude) to draft parts of the code, the benchmark and the text, under the author's direction. The author reviewed the material, ran the reported checks, and takes full responsibility for the content.

References

[1] Apache Seata. Alibaba Seata resolves idempotence, dangling, and empty rollback issues in TCC mode. `https://seata.apache.org/blog/seata-tcc-fence/`, 2026. Describes the TCC fence (table tcc_fence_log) introduced in Seata 1.5.1; undated page, accessed 2026-10-21.

[2] Arm Limited. *Arm Architecture Reference Manual for A-profile Architecture*. Arm Limited, issue l.b edition, 2025. URL `https://developer.arm.com/documentation/ddi0487/latest`.

[3] Arm Limited. *AMBA AXI Protocol Specification*. Arm Limited, issue l edition, 2025. URL `https://developer.arm.com/documentation/ihi0022/latest`.

[4] Joe Armstrong. *Making Reliable Distributed Systems in the Presence of Software Errors*. PhD thesis, The Royal Institute of Technology (KTH), Stockholm, Sweden, 2003.

[5] Jagdish Bansiya, Letha Etzkorn, Carl Davis, and Wei Li. A class cohesion metric for object-oriented designs. *Journal of Object-Oriented Programming*, 11(8):47–52, 1999.

[6] Philip A. Bernstein, Vassos Hadzilacos, and Nathan Goodman. *Concurrency Control and Recovery in Database Systems*. Addison-Wesley, 1987. ISBN 0-201-10715-5.

[7] Gérard Berry and Georges Gonthier. The Esterel synchronous programming language: Design, semantics, implementation. *Science of Computer Programming*, 19(2):87–152, 1992. doi: 10.1016/0167-6423(92)90005-V.

[8] Andrew D. Birrell and Bruce Jay Nelson. Implementing remote procedure calls. *ACM Transactions on Computer Systems*, 2(1):39–59, 1984. doi: 10.1145/2080.357392.

[9] Mike Bishop. HTTP/3. RFC 9114, June 2022. URL `https://www.rfc-editor.org/info/rfc9114`. Section 4.1.1 checked against the text at https://httpwg.org/specs/rfc9114.html.

[10] Roberto Bruni, Hernán C. Melgratti, and Ugo Montanari. Theoretical foundations for compensations in flow composition languages. In *Proceedings of the 32nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '05)*, pages 209–220. ACM, 2005. doi: 10.1145/1040305.1040323.

[11] Glenn Bruns and Patrice Godefroid. Model checking partial state spaces with 3-valued temporal logics. In *Computer Aided Verification (CAV '99)*, volume 1633 of *Lecture Notes in Computer Science*, pages 274–287. Springer, 1999. doi: 10.1007/3-540-48683-6_25.

[12] Sebastian Burckhardt, Chris Gillum, David Justo, Konstantinos Kallas, Connor McMahon, and Christopher S. Meiklejohn. Durable functions: Semantics for stateful serverless. *Proceedings of the ACM on Programming Languages*, 5(OOPSLA):133:1–133:27, 2021. doi: 10.1145/3485510.

[13] Michael Butler, C. A. R. Hoare, and Carla Ferreira. A trace semantics for long-running transactions. In *Communicating Sequential Processes: The First 25 Years*, volume 3525 of *Lecture Notes in Computer Science*, pages 133–150. Springer, 2005. doi: 10.1007/11423348_8.

[14] Sergey Bykov, Alan Geller, Gabriel Kliot, James R. Larus, Ravi Pandya, and Jorgen Thelin. Orleans: Cloud computing for everyone. In *Proceedings of the 2nd ACM Symposium on Cloud Computing (SoCC '11)*. ACM, 2011. doi: 10.1145/2038916.2038932.

[15] Jérémie Chassaing. Functional event sourcing decider. `https://thinkbeforecoding.com/post/2021/12/17/functional-event-sourcing-decider`, December 2021. Blog post, Accessed 2026-10-21.

[16] Shyam R. Chidamber and Chris F. Kemerer. A metrics suite for object oriented design. *IEEE Transactions on Software Engineering*, 20(6):476–493, 1994. doi: 10.1109/32.295895.

[17] Alistair Cockburn. Hexagonal architecture. `https://alistair.cockburn.us/hexagonal-architecture/`, 2005. Accessed 2026-10-21.

[18] Evan Czaplicki. Elm: Concurrent FRP for functional GUIs. Senior thesis, Harvard University, 2012.

[19] Edward B. Eichelberger. Hazard detection in combinational and sequential switching circuits. *IBM Journal of Research and Development*, 9(2):90–99, 1965. doi: 10.1147/rd.92.0090.

[20] Noah Eisen and Eric Gribkoff. gRPC retry design (gRFC A6). `https://github.com/grpc/proposal/blob/master/A6-client-retries.md`, 2024. Last updated 2024-08-23; section "Transparent Retries", accessed 2026-10-21.

[21] E. N. Elnozahy, Lorenzo Alvisi, Yi-Min Wang, and David B. Johnson. A survey of rollback-recovery protocols in message-passing systems. *ACM Computing Surveys*, 34(3):375–408, 2002. doi: 10.1145/568522.568525.

[22] Eric Evans. *Domain-Driven Design: Tackling Complexity in the Heart of Software*. Addison-Wesley, 2003. ISBN 0-321-12521-5.

[23] Roy T. Fielding, Mark Nottingham, and Julian Reschke. HTTP semantics. RFC 9110, STD 97, June 2022. URL `https://www.rfc-editor.org/info/rfc9110`.

[24] Svend Frølund and Rachid Guerraoui. X-ability: A theory of replication. *Distributed Computing*, 14(4): 231–249, 2001. doi: 10.1007/s004460100065.

[25] Hector Garcia-Molina and Kenneth Salem. Sagas. In Umeshwar Dayal and Irving L. Traiger, editors, *Proceedings of the 1987 ACM SIGMOD International Conference on Management of Data, San Francisco, CA, USA, May 27–29, 1987*, pages 249–259. ACM Press, 1987. doi: 10.1145/38713.38742.

[26] Patrice Godefroid. Model checking for programming languages using VeriSoft. In *Proceedings of the 24th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '97)*, pages 174–186. ACM, 1997. doi: 10.1145/263699.263717.

[27] Jim Gray. The transaction concept: Virtues and limitations. In *Proceedings of the Seventh International Conference on Very Large Data Bases (VLDB), Cannes, France, September 1981*, pages 144–154, 1981. Also Tandem Technical Report TR 81.3; pages checked against the VLDB 1981 proceedings index.

[28] Vassos Hadzilacos. A knowledge theoretic analysis of atomic commitment protocols. In *Proceedings of the Sixth ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems (PODS '87)*, pages 129–134. ACM, 1987. doi: 10.1145/28659.28672.

[29] Nicolas Halbwachs, Paul Caspi, Pascal Raymond, and Daniel Pilaud. The synchronous data flow programming language LUSTRE. *Proceedings of the IEEE*, 79(9):1305–1320, 1991. doi: 10.1109/5.97300.

[30] Joseph Y. Halpern and Yoram Moses. Knowledge and common knowledge in a distributed environment. *Journal of the ACM*, 37(3):549–587, 1990. doi: 10.1145/79147.79161.

[31] Pat Helland. Life beyond distributed transactions: an apostate's opinion. In *Third Biennial Conference on Innovative Data Systems Research (CIDR 2007), Asilomar, CA, USA, January 7–10, 2007, Online Proceedings*, pages 132–141. www.cidrdb.org, 2007. URL `http://cidrdb.org/cidr2007/papers/cidr07p15.pdf`.

[32] Pat Helland and David Campbell. Building on quicksand. In *Fourth Biennial Conference on Innovative Data Systems Research (CIDR 2009), Asilomar, CA, USA, January 4–7, 2009, Online Proceedings*. www.cidrdb.org, 2009.

[33] Carl Hewitt, Peter Boehler Bishop, and Richard Steiger. A universal modular ACTOR formalism for artificial intelligence. In Nils J. Nilsson, editor, *Proceedings of the 3rd International Joint Conference on Artificial Intelligence (IJCAI), Stanford, CA, USA, August 20–23, 1973*, pages 235–245. William Kaufmann, 1973. URL `https://www.ijcai.org/Proceedings/73/Papers/027B.pdf`.

[34] Martin Hitz and Behzad Montazeri. Measuring coupling and cohesion in object-oriented systems. In *Proceedings of the International Symposium on Applied Corporate Computing (ISACC '95)*, Monterrey, Mexico, 1995.

[35] Daniel Jackson. *Software Abstractions: Logic, Language, and Analysis*. MIT Press, Cambridge, MA, revised edition, 2012. ISBN 978-0-262-01715-2.

[36] Myron King, Nirav Dave, and Arvind. Automatic generation of hardware/software interfaces. In *Proceedings of the 17th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS XVII)*, pages 325–336. ACM, 2012. doi: 10.1145/2150976.2151011.

[37] Kyle Kingsbury. Jepsen: Distributed systems safety research. `https://jepsen.io`, 2026. Accessed 2026-10-21.

[38] Kyle Kingsbury and Peter Alvaro. Elle: Inferring isolation anomalies from experimental observations. *Proceedings of the VLDB Endowment*, 14(3):268–280, 2020. doi: 10.14778/3430915.3430918.

[39] Stephen Cole Kleene. *Introduction to Metamathematics*. North-Holland, Amsterdam, 1952. Co-published with P. Noordhoff (Groningen) and D. Van Nostrand (New York).

[40] Hermann Kopetz and Günther Bauer. The time-triggered architecture. *Proceedings of the IEEE*, 91(1): 112–126, 2003.

[41] Kim Guldstrand Larsen and Bent Thomsen. A modal process logic. In *Proceedings of the Third Annual Symposium on Logic in Computer Science (LICS '88)*, pages 203–210. IEEE Computer Society, 1988. doi: 10.1109/LICS.1988.5119.

[42] Brandur Leach. Implementing Stripe-like idempotency keys in Postgres. `https://brandur.org/idempotency-keys`, October 2017. Published 2017-10-27; text checked against its source in github.com/brandur/sorg, accessed 2026-10-21.

[43] Jiapeng Li. Where does exactly-once live? Model, harness, and tool-contract effects on duplicate side effects in LLM agents, 2026.

[44] Marten Lohstroh, Christian Menard, Soroush Bateni, and Edward A. Lee. Toward a lingua franca for deterministic concurrent systems. *ACM Transactions on Embedded Computing Systems*, 20(4): 36:1–36:27, 2021. doi: 10.1145/3448128.

[45] Alan MacCormack, John Rusnak, and Carliss Y. Baldwin. Exploring the structure of complex software designs: An empirical study of open source and proprietary code. *Management Science*, 52(7): 1015–1030, 2006. doi: 10.1287/mnsc.1060.0552.

[46] Kiwan Maeng and Brandon Lucia. Supporting peripherals in intermittent systems with just-in-time checkpoints. In *Proceedings of the 40th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '19)*, pages 1101–1116. ACM, 2019. Pages from the NSF Public Access Repository record, https://par.nsf.gov/biblio/10129925.

[47] Robert C. Martin. OO design quality metrics: An analysis of dependencies. White paper, October 1994. Dated October 28, 1994.

[48] Robert C. Martin. *Clean Architecture: A Craftsman's Guide to Software Structure and Design*. Robert C. Martin Series. Prentice Hall, 2017. ISBN 978-0-13-449416-6.

[49] Thomas J. McCabe. A complexity measure. *IEEE Transactions on Software Engineering*, SE-2(4): 308–320, 1976. doi: 10.1109/TSE.1976.233837.

[50] Eagon Meng and Daniel Jackson. What you see is what it does: A structural pattern for legible software. In *Proceedings of the 2025 ACM SIGPLAN International Symposium on New Ideas, New Paradigms, and Reflections on Programming and Software (Onward! '25)*, pages 178–193. ACM, 2025. doi: 10.1145/3759429.3762628.

[51] C. Mohan, Donald J. Haderle, Bruce G. Lindsay, Hamid Pirahesh, and Peter M. Schwarz. ARIES: A transaction recovery method supporting fine-granularity locking and partial rollbacks using write-ahead logging. *ACM Transactions on Database Systems*, 17(1):94–162, 1992. doi: 10.1145/128765.128770.

[52] Suvam Mukherjee, Nitin John Raj, Krishnan Govindraj, Pantazis Deligiannis, Chandramouleswaran Ravichandran, Akash Lal, Aseem Rastogi, and Raja Krishnaswamy. Reliable state machines: A framework for programming reliable cloud services. In *33rd European Conference on Object-Oriented Programming (ECOOP 2019)*, volume 134 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 18:1–18:29. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2019. URL `https://drops.dagstuhl.de/opus/volltexte/2019/10810`.

[53] Narayana project. Narayana transaction manager, source code. `https://github.com/jbosstm/narayana`, 2026. Methods BasicAction.updateHeuristic (ArjunaCore) and SubordinateAtomicAction.doCommit (ArjunaJTA), main branch, accessed 2026-10-21.

[54] Edmund B. Nightingale, Peter M. Chen, and Jason Flinn. Speculative execution in a distributed file system. In *Proceedings of the 20th ACM Symposium on Operating Systems Principles (SOSP '05)*, pages 191–205. ACM, 2005. doi: 10.1145/1095810.1095829.

[55] Edmund B. Nightingale, Kaushik Veeraraghavan, Peter M. Chen, and Jason Flinn. Rethink the sync. In

Proceedings of the 7th USENIX Symposium on Operating Systems Design and Implementation (OSDI '06), pages 1–14. USENIX Association, 2006.

[56] Rishiyur S. Nikhil. Bluespec system verilog: Efficient, correct RTL from high level specifications. In *Proceedings of the 2nd ACM/IEEE International Conference on Formal Methods and Models for Co-Design (MEMOCODE 2004)*, pages 69–70. IEEE, 2004. doi: 10.1109/MEMCOD.2004.1459818.

[57] Guy Pardon and Cesare Pautasso. Atomic distributed transactions: A RESTful design. In *Proceedings of the 23rd International Conference on World Wide Web (WWW '14 Companion)*, pages 943–948. ACM, 2014. doi: 10.1145/2567948.2579221.

[58] David Lorge Parnas. On the criteria to be used in decomposing systems into modules. *Communications of the ACM*, 15(12):1053–1058, 1972. doi: 10.1145/361598.361623.

[59] PCI-SIG. *PCI Express Base Specification, Revision 7.0*. PCI-SIG, 2025.

[60] Dewayne E. Perry and Alexander L. Wolf. Foundations for the study of software architecture. *ACM SIGSOFT Software Engineering Notes*, 17(4):40–52, 1992. doi: 10.1145/141874.141884.

[61] Protocol Buffers authors. descriptor.proto: the `idempotency_level` method option. `https://github.com/protocolbuffers/protobuf/blob/main/src/google/protobuf/descriptor.proto`, 2026. Values IDEMPOTENCY_UNKNOWN, NO_SIDE_EFFECTS and IDEMPOTENT; accessed 2026-10-21.

[62] Ganesan Ramalingam and Kapil Vaswani. Fault tolerance via idempotence. In *Proceedings of the 40th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL '13)*, pages 249–262. ACM, 2013. doi: 10.1145/2429069.2429100.

[63] Chris Richardson. *Microservices Patterns: With Examples in Java*. Manning, 2018. ISBN 978-1-61729-454-9.

[64] Jerome H. Saltzer, David P. Reed, and David D. Clark. End-to-end arguments in system design. *ACM Transactions on Computer Systems*, 2(4):277–288, 1984. doi: 10.1145/357401.357402.

[65] Mary Shaw and David Garlan. *Software Architecture: Perspectives on an Emerging Discipline*. Prentice Hall, Upper Saddle River, NJ, 1996. ISBN 0-13-182957-2.

[66] James E. Smith and Andrew R. Pleszkun. Implementation of precise interrupts in pipelined processors. In *Proceedings of the 12th Annual Symposium on Computer Architecture (ISCA '85), Boston, MA, USA, June 1985*, pages 36–44. IEEE Computer Society, 1985. doi: 10.1145/327010.327125.

[67] Smithy authors. Smithy 2.0 specification: Behavior traits. `https://smithy.io/2.0/spec/behavior-traits.html`, 2026. The readonly, idempotent and idempotencyToken traits; checked against docs/source-2.0/spec/behavior-traits.rst in github.com/smithy-lang/smithy, accessed 2026-10-21.

[68] Robert E. Strom and Shaula Yemini. Optimistic recovery in distributed systems. *ACM Transactions on Computer Systems*, 3(3):204–226, 1985. doi: 10.1145/3959.3962.

[69] Milijana Surbatovich, Limin Jia, and Brandon Lucia. I/O dependent idempotence bugs in intermittent systems. *Proceedings of the ACM on Programming Languages*, 3(OOPSLA), 2019. doi: 10.1145/3360609.

[70] Temporal Technologies. Temporal events reference. `https://docs.temporal.io/references/events`, 2026. ActivityTaskScheduled, ActivityTaskStarted (recorded when the task is dispatched to a worker) and ActivityTaskTimedOut (a timeout failure that holds the timeout type); accessed 2026-10-21 via github.com/temporalio/documentation (events.mdx).

[71] Temporal Technologies. What is a Temporal retry policy? `https://docs.temporal.io/encyclopedia/retry-policies`, 2026. Activities retry by default, with unlimited maximum attempts; accessed 2026-10-21 via github.com/temporalio/documentation.

[72] Temporal Technologies. Detecting activity failures. `https://docs.temporal.io/encyclopedia/detecting-activity-failures`, 2026. Schedule-To-Start (until a worker picks up the task; never retried) and Start-To-Close timeouts; accessed 2026-10-21 via github.com/temporalio/documentation (detecting-activity-failures.mdx).

[73] Martin Thomson and Cory Benfield. HTTP/2. RFC 9113, June 2022. URL `https://www.rfc-editor.org/info/rfc9113`. Section 8.7 checked against the text at https://httpwg.org/specs/rfc9113.html.

[74] Vladislav Tomilov. Ballarch: Core 2.0 specification, reference implementations and benchmark. `https://github.com/4wl2d/Pokeball`, 2026. Version 2.0.0-rc.1. The repository keeps the project's earlier name, Pokeball, until it is renamed to Ballarch.

[75] Artem Trofimov and Boris Novikov. When tool calls succeed but workflows fail: Anomalies at the agent-tool boundary, 2026.

[76] Bas C. van Fraassen. Singular terms, truth-value gaps, and free logic. *The Journal of Philosophy*, 63 (17):481–495, 1966. doi: 10.2307/2024549.

[77] X/Open Company Ltd. *Distributed Transaction Processing: The XA Specification*. X/Open, 1991.

[78] Eren Yildiz, Saad Ahmed, Bashima Islam, Josiah Hester, and Kasim Sinan Yildirim. Efficient and safe I/O operations for intermittent systems. In *Proceedings of the Eighteenth European Conference on Computer Systems (EuroSys '23)*, pages 63–78. ACM, 2023. doi: 10.1145/3552326.3587435.

[79] Terry Yin. Lizard: An extensible cyclomatic complexity analyzer. `https://github.com/terryyin/lizard`, 2026. Version 1.24.1, accessed 2026-10-21.